

Short Course in Artificial Intelligence

19th June 2024

Lecture 2: AI Applications


PNC 2024

46


46

Traditional Programming and Machine Learning

Traditional Programming



Machine Learning



PNC 2024

47

47

Traditional Programming and Machine Learning

◆ Traditional Programming

◆ Machine Learning

Find the square of a number

Write Python code:

```
>>> x = int(input())
```

```
2
```

```
>>> print('the square is', x*x)
the square is 4
```

Provide examples:

```
[ [2,4],[3,9],[4,16],[5,25]...]
```

Train a ML model

Apply the **model** to new data

Traditional Programming and Machine Learning


Why is **Machine Learning** so **useful** ?

Because we often have problems for which we do not know a **formula**

Example: Is an image a **cat** or **dog** ?




cat or dog ?



Data as Vectors

- ◆ Every type of data can be represented as a vector
- ◆ Image: ($w \times h \times 3$) for RGB encoding



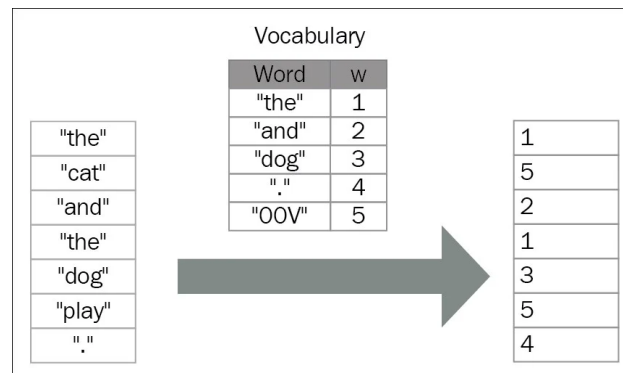
PNC 2024

50

50

Data as Vectors

- ◆ Every type of data can be represented as a vector
- ◆ Words and text:



PNC 2024

51

51

Data as Vectors


- ◆ Every type of data can be represented as a vector
- ◆ Data:
 - Yes – No (1) – (0)
 - User profile (Age, Income, family...)
 - ...

Artificial Intelligence – Machine Learning

- ◆ Learning from **data**
 - Data = set of objects $\{x\}$ $\{images\}$
- 1. Unsupervised learning
 - Groups the objects by similarity
- 2. Supervised learning $\{(image_1, cat), (image_2, dog)\}$
 - Objects have labels $y: \{(x, y)\}$
 - From training data, build a model to predict the label
 $y = f(x)$
- 3. Learning by reinforcement
 - Learn based on reward *(like Tic-Tac-Toe)*

Reinforcement Learning


- ◆ Data is **action**, action get **reward**



- ◆ Learning by adjusting **rewards**
- ◆ Typical example: games, moves are **rewarded** based on **winning** or **losing** game

Unsupervised Machine Learning


- ◆ Data = set of objects $\{x\}$
- ◆ Cluster similar objects together into groups



- ◆ Typical application: marketing, group customers who are likely to buy the same products

Supervised Machine Learning


- ◆ Data = set of objects with label $\{(x, y)\}$
- ◆ Find a model to guess the label y from the value x




- ◆ Typical application: prediction of a value

Machine Learning

◆ Unsupervised



◆ Supervised



Supervised Machine Learning

◆ Which data ?

- Image → Computer Vision CV
- Text → Natural Language Processing NLP
- Data → Example: shopping user data

◆ Which applications ?

- Retail: predict what customers will buy
- Finance: predict future stock value
- Health: automatic diagnostic
- Spam detection: fraudulent emails or sms
- Weather forecast: predict future rain, storm, typhoon
- Vision: face recognition
- Natural Language: chatbots
- ...

Supervised Machine Learning

◆ Data:

- Object: $x = (x_1, x_2, \dots, x_n)$
- Training data: set $\{x_i\} = \{(x_{i1}, x_{i2}, \dots, x_{in})\}$
- Each object x_i has a label y_i :
 $\{(x_i, y_i)\} = \{(x_{i1}, x_{i2}, \dots, x_{in}, y_i)\}$

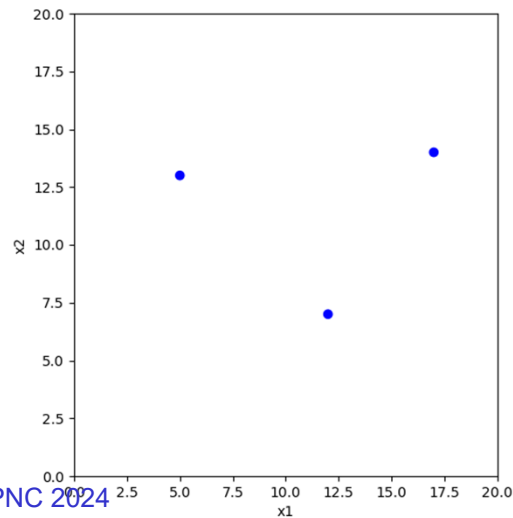
◆ Example: student s has grade x at course C

- We predict grade at C_3 from grades at (C_1, C_2)

Student	C3	C1	C2
s1	16	17	14
s2	9	12	7
s3	7	5	13
...			

Supervised Machine Learning

◆ Geometrical representation



Student	$x_1=C_1$	$x_2=C_2$
s1	17	14
s2	12	7
s3	5	13
...		

PNC 2024

60

60

Supervised Machine Learning

◆ To make it simpler, we only predict if the student **succeeds** or **fails** in C_3

◆ Data:

Student	C3	C3	C1	C2
s1	S	16	17	14
s2	F	9	12	7
s3	F	7	5	13
...				

◆ For easier mathematics, we note:

- Success: S $y_1 = +1$
- Failure : F $y_2 = -1$ $y_3 = -1$


PNC 2024

61

61

Supervised Machine Learning

◆ Geometrical representation



Student	C3	x1=C1	x2=C2
s1	S	17	14
s2	F	12	7
s3	F	5	13
...			


PNC 2024

62

62

Supervised Machine Learning

◆ Linear Classifier



Student	C3	x1=C1	x2=C2
s1	S	17	14
s2	F	12	7
s3	F	5	13
...			


PNC 2024

63

63

Supervised Machine Learning

◆ Linear Classifier: $w_1x_1 + w_2x_2 + w_0 = 0$



w_1	0
w_2	1
w_0	-10
$x_2 - 10 = 0$	

w_1	3
w_2	4
w_0	-80
$3x_1 + 4x_2 - 80 = 0$	

Supervised Machine Learning


◆ The boundary of the linear classifier is the line:

$$w_1x_1 + w_2x_2 + w_0 = 0$$

◆ The line splits the space into two half-spaces

◆ $w_1x_1 + w_2x_2 + w_0 > 0$

◆ $w_1x_1 + w_2x_2 + w_0 < 0$



Supervised Machine Learning

- ◆ How to use a linear classifier ?
- ◆ Given a set of points, we define an **error function** and we find the values of (w_1, w_2, w_0) that minimize this function
 - The common minimization technique is called “Gradient Descent”
- ◆ With the values of (w_1, w_2, w_0) we can classify any point (x_1, x_2) :
 - we compute $w_1x_1 + w_2x_2 + w_0$
 - if $w_1x_1 + w_2x_2 + w_0 > 0$, the classifier predicts $y = +1$
 - if $w_1x_1 + w_2x_2 + w_0 < 0$, the classifier predicts $y = -1$


PNC 2024

66

66

Supervised Machine Learning

- ◆ Linear Classifier: $w_1x_1 + w_2x_2 + w_0 = 0$



$sign(x_2 - 10)$	
(17,14)	+4
(12,7)	-3
(5,13)	+3
1 error	

$sign(3x_1 + 4x_2 - 80)$	
(17,14)	27
(12,7)	-16
(5,13)	-13
0 error	


PNC 2024

67

67

Intuition for Gradient Descent

- ◆ We want to minimize an Error function $E(w)$
Should we increase or decrease w to reduce the value of $E(w)$?




PNC 2024

68

68

Intuition for Gradient Descent

- ◆ We use the derivative to reduce the value $E(w)$
 - if the derivative is negative, we increase $w \rightarrow w + \lambda$
 - if the derivative is positive, we decrease $w \rightarrow w - \lambda$



PNC 2024


69

69

Intuition for Gradient Descent

◆ Problems:

- if λ is too small, it takes very long to reach \bar{w}
- if λ is too big, we miss \bar{w} and we need to come back



Intuition for Gradient Descent

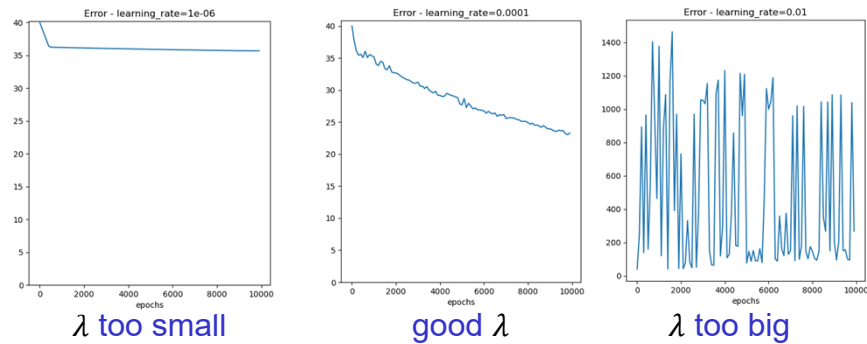
◆ Gradient Descent procedure:

- start with an initial value of w
- compute the derivative
- use it to modify the value of w by an amount that depends on the learning rate λ
- check the error to make sure that λ is not too big
- iterate until convergence
 - the iterations are called “epoch”

Intuition for Gradient Descent

◆ Gradient Descent procedure:

- to find a good learning rate, watch how the value of the error changes with the number of epochs



PNC 2024

72

72

Intuition for Gradient Descent

- ◆ A good value for the learning rate may be different for each function
- ◆ There are many algorithms to change the value of the learning rate during the epochs:
 - Start with a high learning rate → move fast
 - Later reduce the learning rate → be accurate

PNC 2024

73

73


AI Applications

◆ How to build a AI application ?


◆ Steps:

1. **Collect** data
2. **Prepare** the data (into suitable format)
 - cleaning, labelling, ...
3. **Choose** a model
4. **Train** the model
 - find the best values of the parameters
5. **Evaluate** the model
 - on data that was not used to train
6. **Make predictions**
 - on new data

AI Applications



AI Applications



PNC 2024


76

76

AI Applications

- ◆ Image → Data
- ◆ Face Recognition
- ◆ Emotion detection

- ◆ Self-driving cars



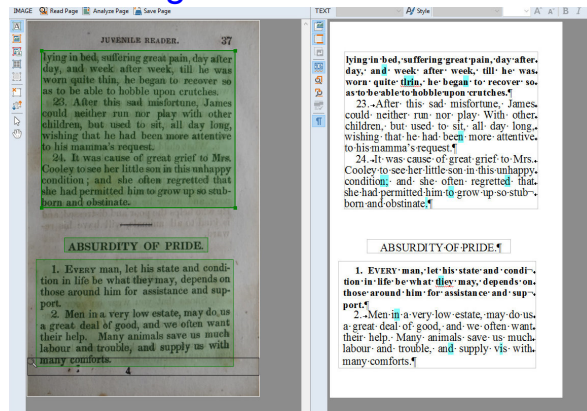
PNC 2024

77

77

AI Applications

- ◆ Image → Text
- ◆ Character Recognition




PNC 2024

78

78

AI Applications

- ◆ Text → Text
- ◆ Automatic Translation



- ◆ Summarization
- ◆ Information search

PNC 2024

79


79

AI Applications

- ◆ Speech Data → Text
- ◆ Speech Recognition
- ◆ Language Learning



- ◆ Text → Speech
- ◆ Speech Synthesis



PNC 2024

80

80

AI Applications

- ◆ Marketing: targeted advertising
- ◆ Is the customer **X** willing to buy product **Y**
- ◆ Data:
 - customer age
 - salary
 - family
 - address
 - products already bought
 - web pages visited
 - ...
- ◆ Prediction: the customer will buy **Y** ? (yes/no)



PNC 2024

81

81

AI Applications

- ◆ Common question to AI engineers from users:
 - how much data do you need ?
- ◆ Common answer:
 - as much as possible
- ◆ But **quality** of data is important

DATA = \$\$
Personal Data = \$\$\$

Quick Introduction to Deep Learning and Demo


*if you are afraid of **formulas**, don't look
at them, just look at the figures*

A Quick Course On Deep Networks

- ◆ Deep Learning is a particular type of models based on the **combination of linear classifiers**
- ◆ Those models are the basis for most of the recent progress in AI/ML
- ◆ They can be applied to (almost) any type of data (text, image, speech, data...), and can be extremely efficient for very diverse tasks.
- ◆ The models can be huge (billions of parameters) and be trained on huge training sets.


A Quick Course On Deep Networks

- ◆ The real plane \mathbb{R}^2
- ◆ Points (vectors) (x_1, x_2)
- ◆ Lines $w_1x_1 + w_2x_2 + w_0 = 0$



A Quick Course On Deep Networks


- ◆ Lines $w_1x_1 + w_2x_2 + w_0 = 0$
- ◆ Half-spaces
 - ◆ $w_1x_1 + w_2x_2 + w_0 > 0$
 - ◆ $w_1x_1 + w_2x_2 + w_0 < 0$



A Quick Course On Deep Networks

- ◆ The ReLU Function (Rectified Linear Unit)
 $ReLU(t) = \max(t, 0)$


- ◆ $ReLU(t) = \begin{cases} t & \text{if } t \geq 0 \\ 0 & \text{if } t < 0 \end{cases}$



- ◆ ReLU on the real plane \mathbb{R}^2
 $ReLU(x_1, x_2) = ReLU(w_1x_1 + w_2x_2 + w_0)$
 $= \max(w_1x_1 + w_2x_2 + w_0, 0)$

A Quick Course On Deep Networks

- ◆ ReLU on the real plane \mathbb{R}^2 : $\text{ReLU}_{w_0, w_1, w_2}(x_1, x_2)$




PNC 2024

88

88

A Quick Course On Deep Networks

- ◆ Combining linear classifiers




PNC 2024

89


89

A Quick Course On Deep Networks

◆ From Linear Classifier



◆ To Deep Networks

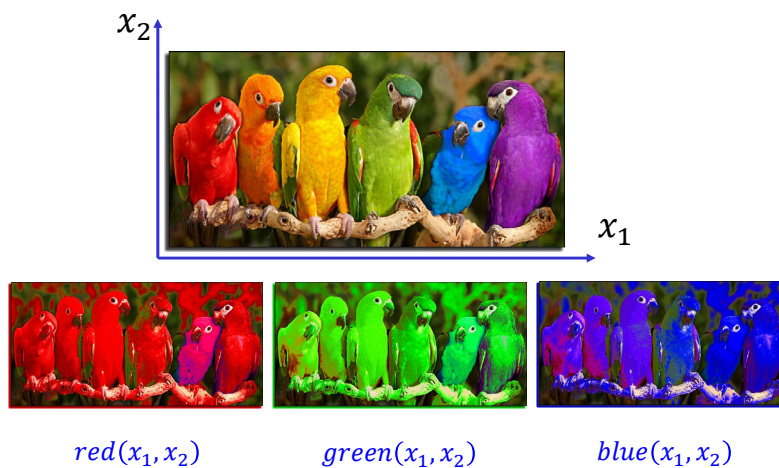


PNC 2024

90

90

A Quick Deep Network Demo




PNC 2024

91

91

A Quick Deep Network



PNC 2024

92

92

Demo: Learning Deep Networks

- ◆ You are given a Javascript code that trains a Deep Network to represent a color image
- ◆ You can see how the network evolves during training
- ◆ You can modify the **number of hidden layers** (just cut and paste one line) to see how deeper networks can better represent the image

PNC 2024

93

93

Demo: Learning Deep Networks

◆ Data :

- One image 200x200




◆ Find models of the color functions:

- $red(x_1, x_2)$
- $green(x_1, x_2)$
- $blue(x_1, x_2)$



Demo: Learning Deep Networks



Demo: Learning Deep Networks

ConvnetJS demo: Image "Painting"

Credits to [Andrei Karpathy](#)


```
layer_defs = []  
// input layer  
layer_defs.push({type:'input', out_sz:1, out_sy:1, out_depth:2}); // 2 inputs: x, y  
// fully connected hidden layer(s)  
layer_defs.push({type:'fc', num_neurons:20, activation:'relu'});  
// output layer  
layer_defs.push({type:'regression', num_neurons:3}); // 3 outputs: r,g,b  
net = new convnetjs.Net();  
net.makeLayers(layer_defs);  
trainer = new convnetjs.SGDTrainer(net, {learning_rate:0.01, momentum:0.9, batch_size:5, l2_decay:0.0});
```

input layer (x_1, x_2)

hidden layer
20 neurons

output layer (r, g, b)

reload network Choose your own image: No file chosen



training image

model image

Learning rate: 0.01

learning rate

The learning rate should probably be decreased over time (slide left) to let the network better overfit the training data. It's nice to not have to worry about overfitting.

PNC 2024

96

96

Demo: Learning Deep Networks


ConvnetJS demo: Image "Painting"

Credits to [Andrei Karpathy](#)

```
layer_defs = []  
// input layer  
layer_defs.push({type:'input', out_sz:1, out_sy:1, out_depth:2}); // 2 inputs: x, y  
// fully connected hidden layer(s)  
layer_defs.push({type:'fc', num_neurons:20, activation:'relu'});  
// output layer  
layer_defs.push({type:'regression', num_neurons:3}); // 3 outputs: r,g,b  
net = new convnetjs.Net();  
net.makeLayers(layer_defs);  
trainer = new convnetjs.SGDTrainer(net, {learning_rate:0.01, momentum:0.9, batch_size:5, l2_decay:0.0});
```

hidden layer
20 neurons

Duplicating this line will add new hidden layers to the model



Learning rate: 0.01

The learning rate should probably be decreased over time (slide left) to let the network better overfit the training data. It's nice to not have to worry about overfitting.

PNC 2024

97

97

Practice 2: Learning Linear Classifier

- ◆ You are provided with a **Python notebook** that already contains the code to train and visualize
- ◆ There are 3 datasets, from **easy** to more **complex**
- ◆ You should experiment with several **learning rates** and the **number of epochs** to achieve a good classifier
- ◆ Note: not all values for the `learning_rate` will lead to a successful training

Practice 2: Learning Linear Classifier

- ◆ Log on <https://colab.research.google.com/>
- ◆ **File** -> **upload notebook**
- ◆ upload `practice2.ipynb`
- ◆ **Files** -> **upload**
- ◆ upload `practice2.data1.data`
- ◆ **Runtime** -> **Run All**
- ◆ you can try different values of the **learning_rate**
- ◆ repeat with `practice2.data2.data` and `practice2.data3.data`

Practice 2: Learning Linear Classifier

- ◆ Data : student and grades

Student	Grade1	Grade2	Success
S1	17	14	1
S2	12	7	-1
S3	5	13	-1

- ◆ Goal: build linear classifier to predict success

